

CSc 361: Computer Communications and Networks (Summer 2016)

Assignment 3: Analysis of IP Protocol

Spec Out: June 28, 2016
Due: 3:30 pm July 22, 2016

1 Goal

The purpose of this assignment is to learn about the IP protocol. You are required to write a C program with the pcap library to analyze a trace of IP datagrams.

2 Introduction

In this assignment, we will investigate the IP protocol, focusing on the IP datagram. We'll do so by analyzing a trace of IP datagrams sent and received by an execution of the *traceroute* program. We'll investigate the various fields in the IP datagram, and study IP fragmentation in detail.

A background of the *traceroute* program is summarized as follows. The *traceroute* program operates by first sending one or more datagrams with the time-to-live (TTL) field in the IP header set to 1; it then sends a series of one or more datagrams towards the same destination with a TTL value of 2; it then sends a series of datagrams towards the same destination with a TTL value of 3; and so on. Recall that a router must decrement the TTL in each received datagram by 1 (actually, RFC 791 says that the router must decrement the TTL by at least one). If the TTL reaches 0, the router returns an ICMP message (type 11 TTL-exceeded) to the sending host. As a result of this behavior, a datagram with a TTL of 1 (sent by the host executing *traceroute*) will cause the router one hop away from the sender to send an ICMP TTL-exceeded message back to the sender; the datagram sent with a TTL of 2 will cause the router two hops away to send an ICMP message back to the sender; the datagram sent with a TTL of 3 will cause the router three hops away to send an ICMP message back to the sender; and so on. In this manner, the host executing *traceroute* can learn the identities of the routers between itself and a chosen destination by looking at the source IP addresses in the datagrams containing the ICMP TTL-exceeded messages. You will be provided with a trace file created by *traceroute*.

Of course, you can create a trace file by yourself. Note that when you create the trace file, you need to use different datagram sizes (e.g., 2500 bytes) so that the captured trace file includes information on fragmentation.

3 Requirements

There are two requirements for this assignment:

3.1 Requirement 1 (R1)

You are required to write a C program with the pcap library to analyze the trace of IP datagrams by an execution of *traceroute*. To make terminologies consistent, in this assignment we call the *source node* as the computer that executes *traceroute*. The *ultimate destination node* refers to the host that is the ultimate destination defined when running *traceroute*. For example, the ultimate destination node is “mit.edu” when you run

```
%traceroute mit.edu 2000
```

In addition, an *intermediate destination node* refers to the router that is not the ultimate destination node but sends back a ICMP message to the source node.

As another example, you can set “don’t fragment” bit and set the number of probes per “ttl” to 5 queries using the following command:

```
%traceroute -F -q 5 mit.edu 200
```

Your program needs to output the following information:

- List the IP address of the source node, the IP address of ultimate destination node, the IP address(es) of the intermediate destination node(s). If multiple the intermediate destination nodes exist, they should be ordered by their hop count to the source node in the increasing order.
- Check the IP header of all datagrams in the trace file, and list the set of values in the *protocol* field of the IP headers. Note that only different values should be listed in a set.
- How many fragments were created from the original datagram? Note that 0 means no fragmentation. Print out the offset (in terms of bytes) of the last fragment of the fragmented IP datagram. Note that if the datagram is not fragmented, the offset is 0.
- Calculate the average and standard deviation of round trip time(s) between the source node and the intermediate destination node (s) and the average round trip time between the source node and the ultimate destination node. The average and the average and standard deviation are calculated over all fragments sent/received between the source nodes and the (intermediate/ ultimate) destination node.

The output format is as follows: (Note that the values do not correspond to any trace file).

The IP address of the source node: 192.168.1.12

The IP address of ultimate destination node: 10.216.216.2

The IP addresses of the intermediate destination nodes:

router 1: 24.218.01.102,

router 2: 24.221.10.103,

router 3: 10.215.118.1.

The values in the protocol field of IP headers:

1: ICMP

17: UDP

73 The number of fragments created from the original datagram is: 3
 74 The offset of the last fragment is: 3680
 75
 76 The avg RRT between 192.168.1.12 and 24.218.01.102 is: 50 ms, the s.d. is: 5 ms
 77 The avg RRT between 192.168.1.12 and 24.221.10.103 is: 100 ms, the s.d. is: 6 ms
 78 The avg RRT between 192.168.1.12 and 10.215.118.1 is: 150 ms, the s.d. is: 5 ms
 79 The avg RRT between 192.168.1.12 and 10.216.216.2 is: 200 ms, the s.d. is: 15 ms
 80

81 3.2 Requirement 2 (R2)

82 From a given set of five *traceroute* trace files, all with the same destination address,

- 83 • determine the number of probes per “ttl” used in each trace file,
- 84 • determine whether or not the sequence of intermediate routers is the same in different trace
85 files,
- 86 • if the sequence of intermediate routers is different in the five trace files, list the difference and
87 explain why,
- 88 • if the sequence of intermediate routers is the same in the five trace files, draw a table as shown
89 below (**warning:** the values in the table do not correspond to any trace files) to compare
90 the RRTs of different traceroute attempts. From the result, which hop is likely to incur the
91 maximum delay? Explain your conclusion.

TTL	Average RRT in trace 1	Average RRT in trace 2	Average RRT in trace 3	Average RRT in trace 4	Average RRT in trace 5
1	0.5	0.7	0.8	0.7	0.9
2	0.9	1	1.2	1.2	1.3
3	1.5	1.5	1.5	1.5	2.5
4	2.5	2	2	2.5	3
5	3	2.5	3	3.5	3.5
6	5	4	5	4.5	4

93 4 Miscellaneous

94 **Important! Please read!**

- 95 • Some intermediate router may only send back one “ICMP TTL exceeded” message for multiple
96 fragments of the same datagram. In this case, please use this ICMP message to calculate RTT
97 for all fragments. For example, Assume that the source sends Frag 1, Frag2, Frag 3 (of the
98 same datagram, ID: 3000). The timestamps for Frag1, Frag2, Frag3 are t_1, t_2, t_3 , respectively.
99 Later, the source receives one “ICMP TTL exceeded” message (ID: 3000). The timestamp is
100 T . Then the RRTs are calculated as: $T - t_1, T - t_2, T - t_3$.
- 101 • More explanation about the output format

The number of fragments created from the original datagram is:

The offset of the last fragment is:

If there are multiple fragmented datagrams, you need to output the above information for each datagram. For example, assume that the source sends two datagrams: D_1 , D_2 , where D_1 and D_2 are the identification of the two datagram. Assume that D_1 has three fragments and D_2 has two fragments. Then output should be:

The number of fragments created from the original datagram D1 is: 3

The offset of the last fragment is: xxx.

The number of fragments created from the original datagram D2 is: 2

The offset of the last fragment is: yyy.

where *xxx* and *yyy* denote the actual number calculated by your program.

- If the tracefile is captured in Linux, the ID of the original UDP cannot be used to match against the ID field within the data of the returned ICMP error message. Nevertheless, the source port number included in the original UDP can be used to match against the ICMP error message.

Note that we do not have such a problem in the tracefile captured in Windows.

5 Deliverables and Marking Scheme

For your final submission of your assignment, you are required to submit your source code to connex. You should include a readme file to tell TA how to compile and run your code. At the last lab session that you attend, you need to demo your assignment to TAs. Nevertheless, before the final due date, you can still make changes on your code and submit a change.txt file to connex to describe the changes after your demo.

The marking scheme is as follows:

Components	Weight
The IP address of the source node (R1)	5
The IP address of ultimate destination node (R1)	5
The IP addresses of the intermediate destination nodes (R1)	10
The correct order of the intermediate destination nodes (R1)	5
The values in the protocol field of IP headers (R1)	5
The number of fragments created from the original datagram (R1)	10
The offset of the last fragment (R1)	10
The avg RRTs (R1)	10
The standard deviations (R1)	5
The number of probes per ttl (R2)	5
Right answer to the second question (R2)	5
Right answer to the third/or fourth question (R2)	10
Make file	5
Code style	5
Readme.txt and change.txt(if any)	5
Total Weight	100

6 Plagiarism

This assignment is to be done individually. You are encouraged to discuss the design of your solution with your classmates, but each person must implement their own assignment.

7 Extra Info: Code Quality

We cannot specify completely the coding style that we would like to see but it includes the following:

1. Proper decomposition of a program into subroutines (and multiple source code files when necessary)—A 500 line program as a single routine won't suffice.
2. Comment—judiciously, but not profusely. Comments also serve to help a marker, in addition to yourself. To further elaborate:
 - (a) Your favorite quote from Star Wars or Douglas Adams' Hitch-hiker's Guide to the Galaxy does not count as comments. In fact, they simply count as anti-comments, and will result in a loss of marks.
 - (b) Comment your code in English. It is the official language of this university.
3. Proper variable names—`leia` is not a good variable name, it never was and never will be.
4. Small number of global variables, if any. Most programs need a very small number of global variables, if any. (If you have a global variable named `temp`, think again.)
5. **The return values from all system calls and function calls listed in the assignment specification should be checked and all values should be dealt with appropriately.**

The End
